THE UNIVERSITY OF TEXAS AT AUSTIN
McCOMBS SCHOOL OF BUSINESS

Diagnostics & Transformations 2

Lecture 13

STA 371G

Salaries of newly hired managers

Salaries of newly hired managers

- Salary (response)
- Manager rating
- Years of experience
- Years since graduation
- Origin (internal or external hire)

Data issues

Data scientists report that they spend 70% of their time on obtaining and cleaning the data. Only 30% is for statistical analysis.

Data issues

Data scientists report that they spend 70% of their time on obtaining and cleaning the data. Only 30% is for statistical analysis.

Never run a regression without exploring and cleaning the data first!

The most common issues:

1. Outliers
2. Missing data
3. Multicollinearity
4. Assumption violations
5. High-leverage points

1. Outliers
2. Missing data
3. Multicollinearity
4. Assumption violations
5. High-leverage points

Exploring the data: Outliers

Boxplots are commonly used to detect outliers. Let's start by looking at the Salary column.

Exploring the data: Outliers

Boxplots are commonly used to detect outliers. Let's start by looking at the Salary column.

```
boxplot(manager$Salary, xlab='Salary', horizontal=T)
```

Exploring the data: Outliers

```
manager[manager$Salary>200,]

    Salary MngrRating YearsExp YrsSinceGrad   Origin
146    511        6.1        2            2 Internal

manager[manager$Salary<0,]

    Salary MngrRating YearsExp YrsSinceGrad   Origin
121    -66        5.7        1            2 Internal
```

Exploring the data: Outliers

```
manager[manager$Salary>200,]

    Salary MngrRating YearsExp YrsSinceGrad   Origin
146    511        6.1        2            2 Internal

manager[manager$Salary<0,]

    Salary MngrRating YearsExp YrsSinceGrad   Origin
121    -66        5.7        1            2 Internal
```

We can deal with outliers in two ways.

- If the result of **errors in the data**, we can try to correct or omit them.

Exploring the data: Outliers

```
manager[manager$Salary>200,]

    Salary MngrRating YearsExp YrsSinceGrad   Origin
146    511        6.1        2            2 Internal

manager[manager$Salary<0,]

    Salary MngrRating YearsExp YrsSinceGrad   Origin
121    -66        5.7        1            2 Internal
```

We can deal with outliers in two ways.

- If the result of **errors in the data**, we can try to correct or omit them.
- If not, omit them from your regression but report on them separately.

Exploring the data: Outliers

Let's omit the outliers by creating a new data set `mclean` that consists of the subset of the data where the salary is between $0 and $200,000.

```
mclean <- manager[manager$Salary>0 &
                  manager$Salary<200,]
```

Exploring the data: Outliers

Let's omit the outliers by creating a new data set `mclean` that consists of the subset of the data where the salary is between \$0 and \$200,000.

```
mclean <- manager[manager$Salary>0 &
                    manager$Salary<200,]
```

Exploring the data: Outliers

```
boxplot(mclean$YearsExp, xlab='Years of Experience',
  horizontal=T)
```

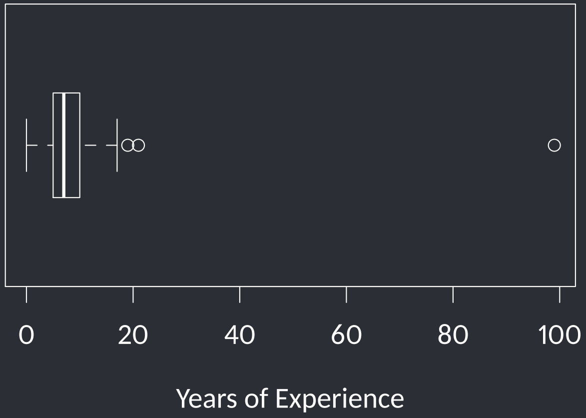

Exploring the data: Outliers

99 must be a code for missing entry in the Years of Experience variable!

Exploring the data: Outliers

99 must be a code for missing entry in the Years of Experience variable!

Let's label all 99s as `NA` (Not Available — R's code for missing data).

Exploring the data: Outliers

99 must be a code for missing entry in the Years of Experience variable!

Let's label all 99s as `NA` (Not Available — R's code for missing data).

```
mclean$YearsExp[mclean$YearsExp == 99] <- NA
```

1. Outliers
2. Missing data
3. Multicollinearity
4. Assumption violations
5. High-leverage points

Exploring the data: Missing entries

Let's see if we have other missing data.

Exploring the data: Missing entries

Let's see if we have other missing data.

```
mclean[!complete.cases(mclean),]

    Salary MngrRating YearsExp YrsSinceGrad   Origin
103     75         NA        8            8 Internal
110     81         NA        9            9 External
124     73        5.9       NA            7 External
154     49        8.0        1            1     <NA>
```

Exploring the data: Missing entries

Let's see if we have other missing data.

```
mclean[!complete.cases(mclean),]
```

```
    Salary MngrRating YearsExp YrsSinceGrad   Origin
103     75         NA        8            8 Internal
110     81         NA        9            9 External
124     73        5.9       NA            7 External
154     49        8.0        1            1     <NA>
```

This isn't surprising—it is very common to have missing entries in your data.

Exploring the data: Missing entries

There are two ways of dealing with missing data:

- Omit the rows that have missing entries in it.

Omiting data is the easiest, but often not the best way, because you lose all the other information available in the same row.

Exploring the data: Missing entries

There are two ways of dealing with missing data:

- Omit the rows that have missing entries in it.
- Try to predict values to fill the missing entries.

Omiting data is the easiest, but often not the best way, because you lose all the other information available in the same row.

Exploring the data: Missing entries

There are two ways of dealing with missing data:

- Omit the rows that have missing entries in it.
- Try to predict values to fill the missing entries.

Omiting data is the easiest, but often not the best way, because you lose all the other information available in the same row.

Exploring the data: Missing entries

There are two ways of dealing with missing data:

- Omit the rows that have missing entries in it.
- Try to predict values to fill the missing entries.

Omiting data is the easiest, but often not the best way, because you lose all the other information available in the same row.

Let's try to fill in some estimates.

Exploring the data: Missing entries

What should we replace the `NA`s in the Manager Rating and Years of Experience columns with?

Exploring the data: Missing entries

What should we replace the NAs in the Manager Rating and Years of Experience columns with?

The simplest way would be to use the averages in the respective columns.

```
mclean$MngrRating[is.na(mclean$MngrRating)] <-
  mean(mclean$MngrRating, na.rm=TRUE)

mclean$YearsExp[is.na(mclean$YearsExp)] <-
  mean(mclean$YearsExp, na.rm=TRUE)
```

Exploring the data: Missing entries

What should we replace the NAs in the Manager Rating and Years of Experience columns with?

The simplest way would be to use the averages in the respective columns.

```
mclean$MngrRating[is.na(mclean$MngrRating)] <-
  mean(mclean$MngrRating, na.rm=TRUE)

mclean$YearsExp[is.na(mclean$YearsExp)] <-
  mean(mclean$YearsExp, na.rm=TRUE)
```

A smarter and more advanced way is to predict the missing data from the other data (using regression!).

Exploring the data: Missing entries

What about the missing data for categorical variables?

Exploring the data: Missing entries

What about the missing data for categorical variables? Let's choose the easy way and omit them.

```
mclean <- na.omit(mclean)
```

Exploring the data: Missing entries

What about the missing data for categorical variables? Let's choose the easy way and omit them.

```
mclean <- na.omit(mclean)
```

This removes all the rows that contain missing entries (only the Origin column has missing entries in this case.)

Exploring the data: Missing entries

What about the missing data for categorical variables? Let's choose the easy way and omit them.

```
mclean <- na.omit(mclean)
```

This removes all the rows that contain missing entries (only the Origin column has missing entries in this case.)

We could also predict the missing entries, or treat the missing entries as a seperate level (e.g. "Unknown").

Exploring the data: Missing entries

Important things to consider:

- While dealing with the missing data, we assume that the data is "Missing Completely at Random" (MCAR).

Exploring the data: Missing entries

Important things to consider:

- While dealing with the missing data, we assume that the data is "Missing Completely at Random" (MCAR).
- If this assumption does not hold (e.g. if the missing data mostly belongs to external hires), the model will be biased.

Exploring the data: Missing entries

Important things to consider:

- While dealing with the missing data, we assume that the data is "Missing Completely at Random" (MCAR).
- If this assumption does not hold (e.g. if the missing data mostly belongs to external hires), the model will be biased.
- Making predictions for missing data based on available data reinforces the existing relationships between variables, so impacts the standard error.

Exploring the data: Missing entries

Important things to consider:

- While dealing with the missing data, we assume that the data is “Missing Completely at Random” (MCAR).
- If this assumption does not hold (e.g. if the missing data mostly belongs to external hires), the model will be biased.
- Making predictions for missing data based on available data reinforces the existing relationships between variables, so impacts the standard error.
- If a lot of data is missing (e.g. more than 5%) for a particular variable, you may have to discard the whole column.

1. Outliers
2. Missing data
3. Multicollinearity
4. Assumption violations
5. High-leverage points

Exploring the data: Multicollinearity

Multicollinearity exists whenever 2+ predictors in a regression model are moderately or highly correlated.

Exploring the data: Multicollinearity

Multicollinearity exists whenever 2+ predictors in a regression model are moderately or highly correlated.

- Any conclusions based on the p-values, coefficients, and confidence intervals of the highly correlated variables will be unreliable.

Correlation between the response and the predictors is good, but correlation between the predictors is not!

Exploring the data: Multicollinearity

Multicollinearity exists whenever 2+ predictors in a regression model are moderately or highly correlated.

- Any conclusions based on the p-values, coefficients, and confidence intervals of the highly correlated variables will be unreliable.
- These statistics will not be stable: adding new data or predictors to the model could drastically change them.

Correlation between the response and the predictors is good, but correlation between the predictors is not!

```r
model <- lm(Salary ~ MngrRating + YearsExp + YrsSinceGrad + Origin,
            data=mclean)
summary(model)
```

```

Call:
lm(formula = Salary ~ MngrRating + YearsExp + YrsSinceGrad +
    Origin, data = mclean)

Residuals:
     Min       1Q   Median       3Q      Max
-19.7766  -4.2842  -0.2906   3.3266  28.2773

Coefficients:
               Estimate Std. Error t value Pr(>|t|)
(Intercept)     54.1521     2.6071  20.771  < 2e-16 ***
MngrRating       4.5147     0.3997  11.296  < 2e-16 ***
YearsExp        -1.5262     1.3790  -1.107 0.270203
YrsSinceGrad     0.7692     1.3833   0.556 0.578976
OriginInternal  -4.7314     1.3878  -3.409 0.000838 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.838 on 149 degrees of freedom
Multiple R-squared:  0.6065,Adjusted R-squared:  0.596
F-statistic: 57.42 on 4 and 149 DF,  p-value: < 2.2e-16
```

```
pairs(MngrRating ~ YearsExp + YrsSinceGrad, data=mclean)
```

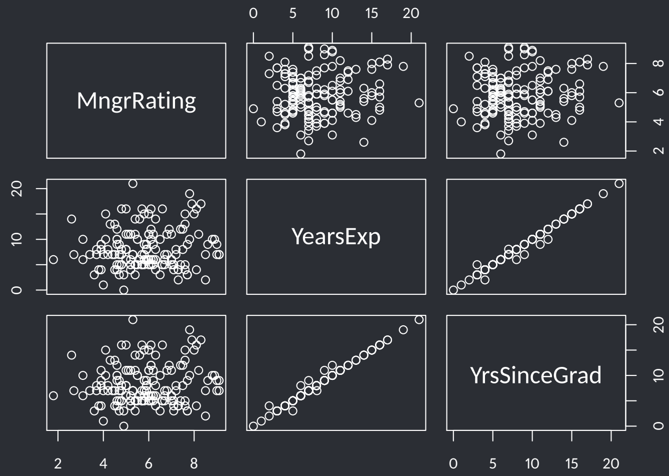

Exploring the data: Multicollinearity

If you check the correlation between the two:

```
cor(mclean$YearsExp, mclean$YrsSinceGrad)

[1] 0.9947616
```

Exploring the data: Multicollinearity

If you check the correlation between the two:

```
cor(mclean$YearsExp, mclean$YrsSinceGrad)
```

```
[1] 0.9947616
```

Any correlation ≥ 0.95 is definitely a problem, but smaller correlations could be problematic too.

Exploring the data: Multicollinearity

A better way to check multicollinearity is using Variance Inflation Factors (VIF):

```
library(car)
vif(model)

  MngrRating     YearsExp YrsSinceGrad       Origin
    1.136002    95.954255    97.011260     1.540448
```

Drop any predictor that has VIF > 5.

Exploring the data: Multicollinearity

A better way to check multicollinearity is using Variance Inflation Factors (VIF):

```
library(car)
vif(model)

  MngrRating     YearsExp YrsSinceGrad       Origin
    1.136002    95.954255    97.011260     1.540448
```

Drop any predictor that has VIF > 5.

Remember: Multicollinearity could exist between more than two predictors (this is why there are only $n - 1$ dummy variables for a categorical variable with n values).

```r
model2 <- lm(Salary ~ MngrRating + YearsExp + Origin, data=mclean)
summary(model2)
```

```
Call:
lm(formula = Salary ~ MngrRating + YearsExp + Origin, data = mclean)

Residuals:
     Min       1Q   Median       3Q      Max 
-19.8115  -4.3474  -0.3964   3.3358  28.1801 

Coefficients:
               Estimate Std. Error t value Pr(>|t|)    
(Intercept)     54.1080     2.5999  20.812  < 2e-16 ***
MngrRating       4.5309     0.3977  11.394  < 2e-16 ***
YearsExp        -0.7651     0.1687  -4.534 1.18e-05 ***
OriginInternal  -4.6467     1.3762  -3.376 0.000935 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.823 on 150 degrees of freedom
Multiple R-squared:  0.6057,Adjusted R-squared:  0.5978 
F-statistic: 76.82 on 3 and 150 DF,  p-value: < 2.2e-16
```

1. Outliers
2. Missing data
3. Multicollinearity
4. Assumption violations
5. High-leverage points

```r
plot(model2)
```

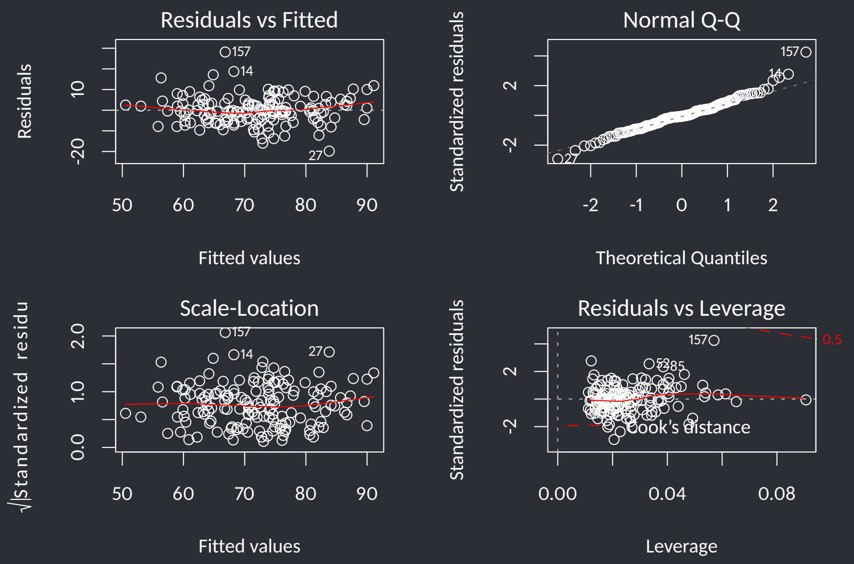

1. Outliers
2. Missing data
3. Multicollinearity
4. Assumption violations
5. High-leverage points

Outliers among the residuals

The data with index 157 seems to be hanging out by itself.

Outliers among the residuals

The data with index 157 seems to be hanging out by itself.

We can display the indices of all of the outliers among the residuals.

```
boxplot(resid(model2))$out

       14        27        52        85       157 
 18.76901 -19.81152  17.14133  15.66079  28.18007 
```

Outliers among the residuals

The data with index 157 seems to be hanging out by itself.

We can display the indices of all of the outliers among the residuals.

```
boxplot(resid(model2))$out
```

```
      14        27        52        85       157 
18.76901 -19.81152  17.14133  15.66079  28.18007 
```

These cases are not predicted well by the model.

Outliers among the residuals

Let's look at row 157:

```
manager[157,]

    Salary MngrRating YearsExp YrsSinceGrad   Origin
157     95          4        1            1 Internal
```

Someone with only 1 year of experience and a poor rating is hired as manager at $95K!

Outliers among the residuals

Let's look at row 157:

```
manager[157,]

    Salary MngrRating YearsExp YrsSinceGrad   Origin
157     95          4        1            1 Internal
```

Someone with only 1 year of experience and a poor rating is hired as manager at $95K!

If you decide that this is an anomaly (e.g. the CEO's son was promoted!) that you don't want to include in your analysis, omit that row and report on it separately in your conclusions.

Influential cases

- The Residuals vs Leverage plot tells about influential cases.

Influential cases

- The Residuals vs Leverage plot tells about influential cases.
- A **high-leverage case** is one that has an unusual combination of predictor values.

Influential cases

- The Residuals vs Leverage plot tells about influential cases.
- A **high-leverage case** is one that has an unusual combination of predictor values.
- An **influential case** is a high-leverage case that also has a high residual: it could change your β values significantly when excluded from your analysis, i.e., it does not follow the overall trend.

Influential cases

- The Residuals vs Leverage plot tells about influential cases.
- A **high-leverage case** is one that has an unusual combination of predictor values.
- An **influential case** is a high-leverage case that also has a high residual: it could change your β values significantly when excluded from your analysis, i.e., it does not follow the overall trend.
- Look for the cases on the upper/lower right corners (beyond the dashed curves).